

Lifting Linear Sketches: Optimal Bounds and Adversarial Robustness

Elena Gribelyuk¹, Honghao Lin², David P. Woodruff², Huacheng Yu¹, Samson Zhou³

¹ Princeton University, ² Carnegie Mellon University, ³ Texas A&M University

Standard Streaming Model

- **Input:** Elements of an underlying frequency vector $x \in \mathbb{Z}^n$, which arrive sequentially one at a time (*worst-case*, fixed in advance).
- **Output:** At the end of the stream, A outputs an approximation of a given function of the stream.
- **Goal:** A should use space *sublinear* in the length m of the input stream and universe size n .

Linear Sketches

- A *linear sketch* is an algorithm that
 1. Samples a sketching matrix $A \in \mathbb{R}^{r \times n}$ and maintains Ax throughout the stream (typically $r \ll n$).
 2. Returns $f(Ax)$ for some estimator f .
- Lower bounds are often proven by selecting a pair of hard distributions \mathcal{D}_1 and \mathcal{D}_2 which exhibit a "gap" for the problem of interest.
- Then, show that $d_{TV}(Ax, Ay)$ is small when A has r rows.

Dimension Lower Bounds

- For many problems, (e.g. operator norm, norm estimation, etc), the hard distributions \mathcal{D}_1 and \mathcal{D}_2 are chosen to be Gaussians (or somewhat "near" Gaussian).
- **Example:** For the problem of estimating $\|x\|_2^2$, pick $\mathcal{D}_1 = \mathcal{N}(0, I_n)$ and $\mathcal{D}_2 = \mathcal{N}(0, (1 + \epsilon)I_n)$. WLOG, A has orthonormal rows. Then, $Ax \sim \mathcal{N}(0, I_r)$, $Ay \sim \mathcal{N}(0, (1 + \epsilon)I_r)$, so A must have $r = \Omega\left(\frac{1}{\epsilon^2} \log 1/\delta\right)$ rows.
- Unfortunately, none of these lower bounds translate to the streaming model!
- **Question:** Is it possible to lift linear sketch lower bounds for continuous inputs to obtain linear sketch lower bounds for discrete inputs?

Adversarially Robust Streaming

- **Input:** Elements of a stream, which arrive sequentially and *adversarially*.
 - **Output:** At each time t , A receives an update u_t , updates its internal state, and returns a *current estimate* r_t , which is recorded by the *adversary*.
- "Future updates may depend on previous estimates"*

Adaptive Attack for Linear Sketches

- Linear sketches for F_p estimation ($p > 0$) are "not robust" to adversarial attacks, i.e. require $\Omega(n)$ dimension [HW13].
- **High-level intuition:** suppose the adversary knows the sketch matrix A : then, a hard distribution is to query $x \in \ker(A)$ or $x = 0^n$, each with probability $1/2$.
- Thus, the adversary will aim to learn the *rowspace* $R(A)$.
- [HW13] gives an adaptive attack which proceeds as follows: initialize $V_1 = \emptyset$.
 1. **Correlation finding:** Find vectors weakly correlated with A orthogonal to V_{i-1} .
 2. **Boosting:** Use these vectors to find strongly correlated vector v .
 3. **Progress:** Set $V_i = \text{span}(V_{i-1}, v)$.
- **Drawback:** All queries are drawn from (continuous) Gaussian distributions with appropriate covariance, and the analysis heavily relies on rotational invariance. This lower bound does not directly translate to the adversarial streaming setting!
- **Question:** Does there exist a sublinear space adversarially robust F_p estimation linear sketch in a finite precision stream?

Main Results

We give a technique for lifting linear sketch lower bounds for continuous inputs to achieve linear sketch lower bounds for discrete inputs!

- **Theorem:** Any adversarially robust streaming algorithm which uses a (finite-precision) linear sketch and B -approximates the F_p moment in a turnstile stream must use $r \geq n - O(\log Bn)$ rows.
- **We also lift linear sketch lower bounds for streaming problems such as operator norm, eigenvalue estimation, compressed sensing, etc.**

Overview of our Approach

Essentially, we want to "simulate" continuous Gaussian queries using discrete Gaussian queries.

- Let $\mathcal{D}_{L,S}$ denote the discrete Gaussian distribution on support L and with covariance matrix $S^T S$.
- Let $x \sim \mathcal{D}_{\mathbb{Z}^n, S}$, $y \sim \mathcal{D}_{A\mathbb{Z}^n, SA^T}$, $g \sim \mathcal{N}(0, AS^T SA^T)$
- **As in continuous case, we want to show $d_{TV}(Ax, y)$ is small on support $A\mathbb{Z}^n$.**
- **Lemma [AR16]: this is true**, under a certain condition for the *orthogonal lattice* to A !
 1. We design a *pre-processing* for the sketching matrix A , which can be applied without loss of generality, and satisfies the above condition.
 2. After applying the pre-processing on sketching matrix A , we show that $Ax + \eta$ and Ag are close in distribution, where η is a uniform noise in the fundamental parallelepiped of the lattice induced by A .
 3. WLOG, assume algorithm sees $Ax + \eta$, since algorithm can always round to recover Ax .